

Real-time Gaming Leaderboard

EDS : +1 point if win the match
each month kick off → new leaderboard
10 users display
5 mill DAU, 25 mill monthly users
10 matches per day
display real-time results

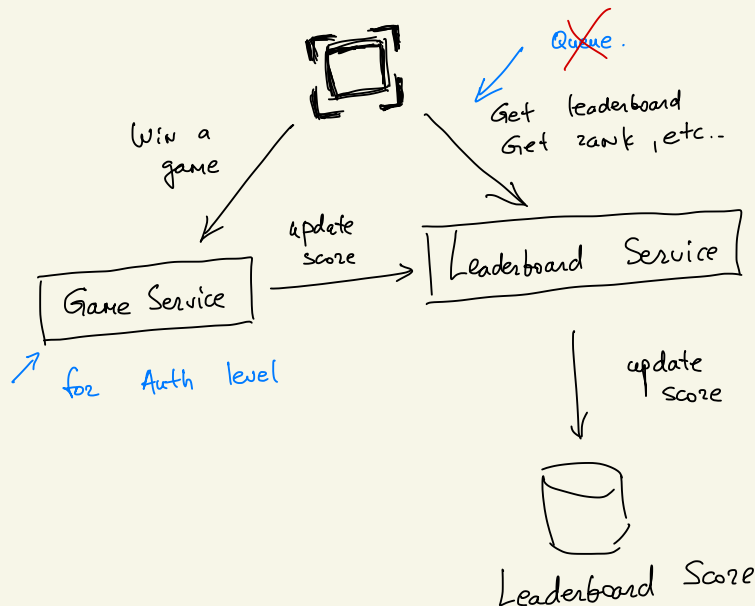
NFR: Real-time update / real-time reflected on the leaderboard
General scalability, availability, reliability

BEE: $\frac{5.000.000 \text{ DAU}}{10^5 \text{ sec}} \approx 50 \text{ QPS}$; Peak: $\text{QPS} \times 5 = 250 \text{ QPS}$
time $\times = 500 \text{ QPS}$
 $+ = 2500 \text{ QPS}$

HLI: API design:

POST /v1/scores update user position on the leaderboard
GET /v1/scores fetch top 10 players
GET /v1/scores/{userId} fetch rank of specific user

HLA:



Database :

RDS - heavy changing data, need to sort
over millions of rows - speed 10s.
No caching for constantly changing data.
Index + Do not performant to scale.
Requires the table scan to determine
the rank.

Redis : in-memory . SortedSets data Type . $O(N)$ for insertion, removal, search.
to improve performance use Binary Search. $O(\log N)$

Implementation using Redis SortedSet :

ZADD - insert user
ZINCRBY - increment the score
ZRANGE / ZREVRANGE - fetch range of users
ZRANK / ZREVRANK - fetch user position

ZINCRBY <key> <increment> <user>

Storage requirement:

name_id	24 char string	} 26 bytes
score	16 bit int (2 bytes)	

25 mill * 26 bytes = 650 MB

CPU, I/O usage:

peak: 250 QPS . This is well within the performance envelope of a single Redis server.

Redis do not support persistence.
Use Redis read replica in configuration.

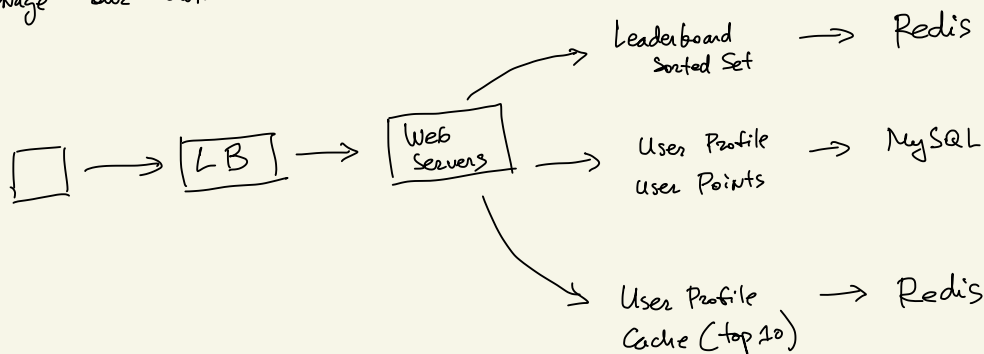
MySQL : user_id
score
timestamp

← data for play history, etc.. services

Additional cache for top 10 players.

- DDD :
1. Cloud ? AWS ?
 2. Scaling Redis.
 3. No SQL ?

Manage our own services :



Using AWS:

API gateway → REST endpoints
Lambda Function

Auto Scaling +

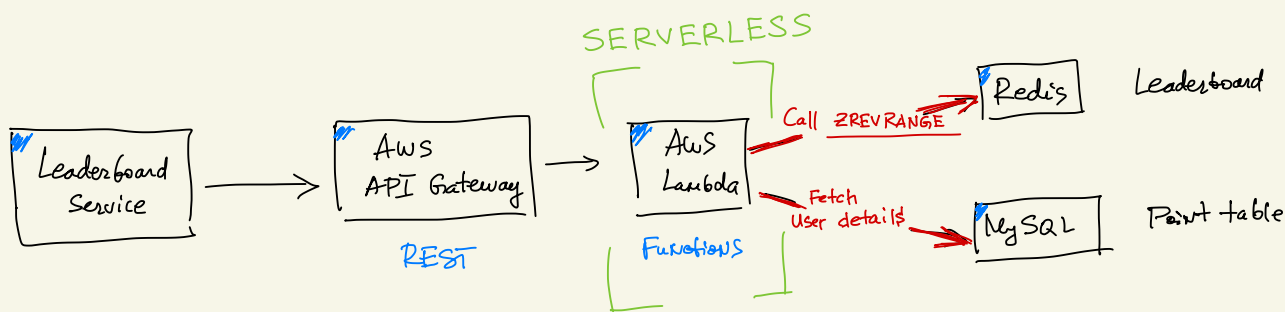
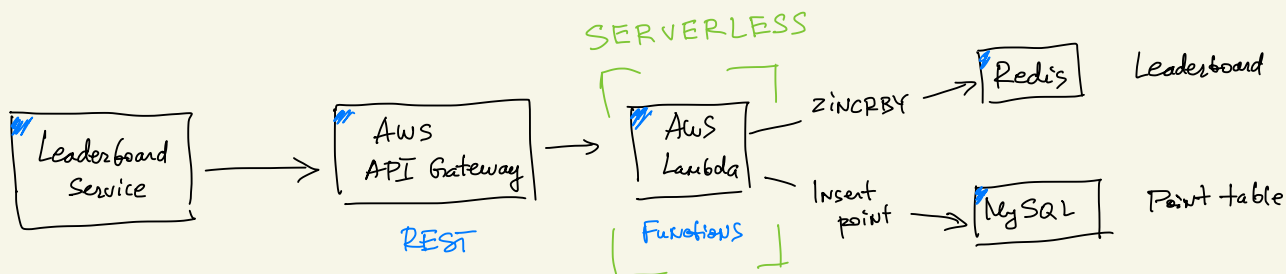
API

GET /v1/scores
GET /v1/scores/{userId}
POST /v1/scores

Lambda Function

Leaderboard Fetch Top 10
Leaderboard Fetch Player Rank
Leaderboard Update Score

invoke
Redis, MySQL
commands



Scaling:

5 mill DAU no need to scale.

↗ × 100 = 500 mill DAU

↗ 65GB × 100 ; ↗ if QPS 250,000 !


Sharding: 1 ... 1000 (points)

Fixed partition:
1..100 , 100..200 , ...
Sorted Set Sorted Set

Will need to fetch top 10 from the last shard
O(1)
900 - 1000 .
Sorted Set

Hash partition:

Redis cluster, hash slots 16384.



Using NoSQL:

1. Optimized for writes
2. Efficiently sort

Amazon DynamoDB, Cassandra, Mongo DB

